

Salary_Analysis_and_Prediction

August 7, 2025

1 Data Science Salary Analysis and Prediction

1.1 1. Import Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
import warnings

warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
sns.set_style('whitegrid')
```

1.2 2. Load and Explore Data

```
[3]: df = pd.read_csv('salaries.csv')
print("Dataset shape:", df.shape)
print("First 5 rows:")
display(df.head())
print("Data types and missing values:")
df.info()
print("Summary statistics:")
display(df.describe(include='all'))
```

Dataset shape: (151445, 11)

First 5 rows:

	work_year	experience_level	employment_type	job_title	salary	\
0	2025	EX	FT	Head of Data	348516	
1	2025	EX	FT	Head of Data	232344	
2	2025	SE	FT	Data Scientist	145400	

3	2025	SE	FT	Data Scientist	81600
4	2025	MI	FT	Engineer	160000

	salary_currency	salary_in_usd	employee_residence	remote_ratio	\
0	USD	348516	US	0	
1	USD	232344	US	0	
2	USD	145400	US	0	
3	USD	81600	US	0	
4	USD	160000	US	100	

	company_location	company_size
0	US	M
1	US	M
2	US	M
3	US	M
4	US	M

Data types and missing values:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 151445 entries, 0 to 151444

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	work_year	151445 non-null	int64
1	experience_level	151445 non-null	object
2	employment_type	151445 non-null	object
3	job_title	151445 non-null	object
4	salary	151445 non-null	int64
5	salary_currency	151445 non-null	object
6	salary_in_usd	151445 non-null	int64
7	employee_residence	151445 non-null	object
8	remote_ratio	151445 non-null	int64
9	company_location	151445 non-null	object
10	company_size	151445 non-null	object

dtypes: int64(4), object(7)

memory usage: 12.7+ MB

Summary statistics:

	work_year	experience_level	employment_type	job_title	\
count	151445.000000	151445	151445	151445	
unique	NaN	4	4	422	
top	NaN	SE	FT	Data Scientist	
freq	NaN	87491	150541	18751	
mean	2024.435313	NaN	NaN	NaN	
std	0.671842	NaN	NaN	NaN	
min	2020.000000	NaN	NaN	NaN	
25%	2024.000000	NaN	NaN	NaN	
50%	2025.000000	NaN	NaN	NaN	
75%	2025.000000	NaN	NaN	NaN	

max	2025.000000	NaN	NaN	NaN
-----	-------------	-----	-----	-----

	salary	salary_currency	salary_in_usd	employee_residence \
count	1.514450e+05	151445	151445.000000	151445
unique	NaN	26	NaN	104
top	NaN	USD	NaN	US
freq	NaN	143173	NaN	135506
mean	1.628380e+05	NaN	157527.458411	NaN
std	2.080124e+05	NaN	74150.772377	NaN
min	1.400000e+04	NaN	15000.000000	NaN
25%	1.060000e+05	NaN	105800.000000	NaN
50%	1.470000e+05	NaN	146100.000000	NaN
75%	1.990000e+05	NaN	198000.000000	NaN
max	3.040000e+07	NaN	800000.000000	NaN

	remote_ratio	company_location	company_size
count	151445.000000	151445	151445
unique	NaN	97	3
top	NaN	US	M
freq	NaN	135569	147302
mean	20.938625	NaN	NaN
std	40.620393	NaN	NaN
min	0.000000	NaN	NaN
25%	0.000000	NaN	NaN
50%	0.000000	NaN	NaN
75%	0.000000	NaN	NaN
max	100.000000	NaN	NaN

1.3 3. Data Cleaning and Feature Engineering

```
[4]: print(f"Shape before dropping duplicates: {df.shape}")
df = df.drop_duplicates()
print(f"Shape after dropping duplicates: {df.shape}")

# Feature Engineering
df['is_remote'] = df['remote_ratio'].apply(lambda x: 1 if x > 50 else 0)
df['is_us'] = df['company_location'].apply(lambda x: 1 if x == 'US' else 0)

def categorize_job(title):
    title = title.lower()
    if 'data scientist' in title:
        return 'Data Scientist'
    elif 'data engineer' in title:
        return 'Data Engineer'
    elif 'machine learning' in title or 'ml' in title:
        return 'Machine Learning'
    elif 'ai' in title:
```

```

        return 'AI'
    elif 'analyst' in title:
        return 'Analyst'
    elif 'manager' in title:
        return 'Manager'
    elif 'engineer' in title:
        return 'Engineer'
    elif 'software' in title:
        return 'Software Engineer'
    elif 'research' in title:
        return 'Research'
    else:
        return 'Other'

df['job_family'] = df['job_title'].apply(categorize_job)

print("Cleaned dataframe sample with new features:")
display(df.head())

```

Shape before dropping duplicates: (151445, 11)

Shape after dropping duplicates: (71913, 11)

Cleaned dataframe sample with new features:


	work_year	experience_level	employment_type	job_title	salary	\
0	2025	EX	FT	Head of Data	348516	
1	2025	EX	FT	Head of Data	232344	
2	2025	SE	FT	Data Scientist	145400	
3	2025	SE	FT	Data Scientist	81600	
4	2025	MI	FT	Engineer	160000	

	salary_currency	salary_in_usd	employee_residence	remote_ratio	\
0	USD	348516	US	0	
1	USD	232344	US	0	
2	USD	145400	US	0	
3	USD	81600	US	0	
4	USD	160000	US	100	

	company_location	company_size	is_remote	is_us	job_family
0	US	M	0	1	Other
1	US	M	0	1	Other
2	US	M	0	1	Data Scientist
3	US	M	0	1	Data Scientist
4	US	M	1	1	Engineer

1.4 4. Exploratory Data Analysis (EDA)

```
[5]: plt.figure(figsize=(12, 6))
sns.histplot(df['salary_in_usd'], bins=50, kde=True)
plt.title('Distribution of Salaries in USD', fontsize=16)
plt.xlabel('Salary in USD', fontsize=14)
plt.ylabel('Count', fontsize=14)
plt.show()
```



1.4.1 Salary by Experience Level

```
[6]: plt.figure(figsize=(12, 6))
order = ['EN', 'MI', 'SE', 'EX']
sns.boxplot(x='experience_level', y='salary_in_usd', data=df, order=order)
plt.title('Salary Distribution by Experience Level', fontsize=16)
plt.xlabel('Experience Level', fontsize=14)
plt.ylabel('Salary in USD', fontsize=14)
plt.xticks([0, 1, 2, 3], ['Entry', 'Mid', 'Senior', 'Executive'])
plt.show()
```




1.4.2 Salary by Job Family

```
[7]: plt.figure(figsize=(14, 8))
order = df.groupby('job_family')['salary_in_usd'].median().
      ↪sort_values(ascending=False).index
sns.boxplot(x='job_family', y='salary_in_usd', data=df, order=order)
plt.title('Salary Distribution by Job Family', fontsize=16)
plt.xlabel('Job Family', fontsize=14)
plt.ylabel('Salary in USD', fontsize=14)
plt.xticks(rotation=45, ha='right')
plt.show()
```



1.4.3 Salary Trend by Year

```
[8]: plt.figure(figsize=(12, 6))
sns.lineplot(x='work_year', y='salary_in_usd', data=df, estimator='median',
             errorbar=None, marker='o')
plt.title('Median Salary Trend by Year', fontsize=16)
plt.xlabel('Year', fontsize=14)
plt.ylabel('Median Salary in USD', fontsize=14)
plt.xticks(df['work_year'].unique())
plt.show()
```



1.5 5. Machine Learning Model to Predict Salary

```
[9]: # Define features and target
features = ['experience_level', 'employment_type', 'job_family',
            ↪ 'remote_ratio', 'company_size', 'is_us']
target = 'salary_in_usd'

X = df[features]
y = df[target]

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
            ↪ random_state=42)

# Define preprocessing steps
categorical_features = ['experience_level', 'employment_type', 'job_family',
            ↪ 'company_size']
numerical_features = ['remote_ratio', 'is_us']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

# Create the pipeline
pipeline = Pipeline(steps=[('preprocessor', preprocessor),
```



```

        ('regressor',
         RandomForestRegressor(random_state=42)))

```

1.5.1 Hyperparameter Tuning with GridSearchCV

```

[10]: param_grid = {
        'regressor__n_estimators': [100, 200],
        'regressor__max_depth': [10, 20],
        'regressor__min_samples_leaf': [2, 4]
    }

    grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='r2', n_jobs=-1,
                               verbose=2)
    grid_search.fit(X_train, y_train)

    print("Best parameters found:", grid_search.best_params_)
    best_model = grid_search.best_estimator_

```

```

Fitting 5 folds for each of 8 candidates, totalling 40 fits
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=100; total time= 14.9s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=100; total time= 15.0s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=100; total time= 15.2s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=100; total time= 16.9s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=100; total time= 14.6s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=200; total time= 29.0s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=200; total time= 30.6s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=200; total time= 29.5s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,
regressor__n_estimators=200; total time= 29.2s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=4,
regressor__n_estimators=100; total time= 13.9s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=4,
regressor__n_estimators=100; total time= 13.6s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=4,
regressor__n_estimators=100; total time= 13.2s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=4,
regressor__n_estimators=100; total time= 13.4s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=4,
regressor__n_estimators=100; total time= 14.6s
[CV] END regressor__max_depth=10, regressor__min_samples_leaf=2,

```

[illegible]

```
regressor__n_estimators=200; total time= 42.2s
[CV] END regressor__max_depth=20, regressor__min_samples_leaf=4,
regressor__n_estimators=200; total time= 35.8s
Best parameters found: {'regressor__max_depth': 20,
'regressor__min_samples_leaf': 4, 'regressor__n_estimators': 200}
```

1.5.2 Model Evaluation


```
[11]: y_pred = best_model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"Root Mean Squared Error (RMSE): ${rmse:,.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.5)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'k--', lw=2)
plt.xlabel('Actual Salary')
plt.ylabel('Predicted Salary')
plt.title('Actual vs. Predicted Salaries')
plt.show()
```

Root Mean Squared Error (RMSE): \$67,995.63

R-squared (R2) Score: 0.24



1.6 6. Conclusion

This analysis provides a comprehensive overview of data science salaries. Key findings include:

- Salaries have been steadily increasing from 2020 to 2025.
- Experience level is a major factor in determining salary, with Executive-level professionals earning significantly more.
- The United States is the highest-paying country for data science roles.
- The 'Machine Learning' and 'Data Scientist' job families command the highest salaries.

The machine learning model can predict salaries with a reasonable R-squared value, indicating that the selected features have predictive power. Further improvements could be made by including more features, trying different models like Gradient Boosting, and gathering more data.